

CS259 Project 2: Hierarchical Performance Model for Convolution Kernels

Executive Summary

This report documents the design and implementation of a **mechanistic hierarchical performance model** for 6 convolution kernel variants (naive, smem_w, unroll, smem_wi, regblock, warp) on NVIDIA TITAN V. The model predicts execution time by analyzing three competing bottlenecks at every level of the memory hierarchy:

$$\text{predicted_time} = \max(s_c \cdot t_{\text{compute}}, t_{\text{L2}}, s_m \cdot t_{\text{DRAM}})$$

The three terms correspond to:

1. **Compute bound** — FLOPs vs peak FP32 throughput (scaled by s_c)
2. **L2 bound** — bytes served from on-chip L2 cache at rated L2 bandwidth
3. **DRAM bound** — bytes transferred from off-chip HBM (scaled by s_m)

The key innovations are:

- **Per-kernel mechanistic DRAM traffic** with an L2-pressure-aware log-linear correction
- **True three-level hierarchical bottleneck**: L2 is an active term in the $\max()$, not just a diagnostic
- **Occupancy and coalescing model** for K3 smem_wi that accurately predicts the performance U-curve over tile sizes

Part 1: From Roofline to Mechanistic Models

1.1 The Original Roofline Baseline

The classic **roofline model** predicts performance by placing kernels on a 2D chart:

- **X-axis**: Arithmetic Intensity (AI) = FLOPs / bytes accessed
- **Y-axis**: Achieved throughput (GFLOPS)
- **Roof**: Two lines forming an envelope — compute roof (horizontal) and memory roof (diagonal)

$$\text{predicted_time_s} = \max\left(\frac{\text{FLOPs}}{\text{peak_FP32}}, \frac{\text{DRAM_bytes}}{\text{peak_BW}}\right)$$

Why roofline alone is insufficient:

- Assumes every byte is read exactly once (theoretical minimum traffic).
- Conv1 naive: roofline predicts ~4 ms; actual is 686 ms → **99.4% MAPE**.
- Cannot distinguish between kernels: naive vs smem_wi both have similar theoretical AI but vastly different performance.

1.2 Why Mechanistic Traffic Estimation?

Instead of theoretical-minimum bytes or purely measured bytes, we:

1. Analyze each kernel's access pattern to build an **analytic DRAM traffic estimate**.
2. Fit a small **L2-pressure correction** that accounts for cache evictions and coalescing gaps.
3. Compute **L2-resident traffic** separately from DRAM traffic (tensor reuse from L2 cache).
4. Model **occupancy and coalescing effects** for the smem_wi tile sweep.

Part 2: Kernel-by-Kernel DRAM Traffic Analysis

All kernels operate on:

$$\text{output}[b, y, x, n] = \sum_{k_y} \sum_{k_x} \sum_i \text{weight}[k_y, k_x, i, n] \times \text{input}[b, y+k_y, x+k_x, i]$$

Key tensors (float32):

Tensor	Shape	Size
Weights	(K_y, K_x, N_i, N_n)	$K_y K_x N_i N_n \times 4$ bytes
Input (padded)	$(B, N_y \{+\} K_y \{-\} 1, N_x \{+\} K_x \{-\} 1, N_i)$	$B(N_y \{+\} 2)(N_x \{+\} 2)N_i \times 4$ bytes
Output	(B, N_y, N_x, N_n)	$B N_y N_x N_n \times 4$ bytes

2.1 K0: Naive

Each thread reads all inputs/weights directly from DRAM per output element:

$$\text{DRAM}_{\text{naive}} = B N_y N_x N_n \cdot K_y K_x N_i \cdot 4 \text{ (input) + weights + output}$$

2.2 K1: smem_w (Weights in Shared Memory)

Weights are loaded cooperatively once per block; input still re-fetched per thread. Analytic estimate is the same as naive (SPAD reuse of weights is not modeled in first-order DRAM count — it appears as L2 reuse when weights fit in L2).

2.3 K2: Unroll

Same access pattern as naive with a $1.25\times$ overhead multiplier on input traffic to account for instruction-scheduling pressure and increased memory-side stalls from unrolled loops.

2.4 K3: smem_wi (Weights + Input in Shared Memory — Key Kernel)

Tiling strategy:

1. Divide output into 16×16 spatial tiles (one block per tile).
2. Divide input channels into `smwi_tile_ni` chunks (default = 8).
3. Per chunk: cooperatively load input patch $(18 \times 18 \times \text{tni})$ and weight slice $(3 \times 3 \times \text{tni})$ into shared memory.
4. Compute all outputs for that chunk from SPAD; repeat.

Analytic DRAM traffic:

$$\text{DRAM}_{\text{input}} = B \cdot N_n \cdot \lceil N_y/16 \rceil \cdot \lceil N_x/16 \rceil \cdot \lceil N_i/\text{tni} \rceil \cdot 18^2 \cdot \text{tni} \cdot 4$$

Note: tni cancels algebraically — the total input DRAM traffic is independent of tile size when coalescing is perfect. The tile size matters through **coalescing efficiency** and **occupancy** (see Part 2b).

Weight traffic — if weights fit in L2 ($K_y K_x N_i N_n \times 4 < \text{L2_size}$), weight re-reads are served from L2, not DRAM:

$$\text{DRAM}_{\text{weight}} = \min \left(\text{syn_bytes}, \text{syn_bytes} \times \frac{B \cdot N_n}{\text{num_SMs}} \right)$$

Conv1 example ($N_y=224, N_i=N_n=64$): weights = 144 KB $\ll$ 4.5 MB L2 $\rightarrow$ loaded once; input patches = 15.4 GB $\rightarrow$ exceeds L2 significantly.

Conv2 example ($N_y=14, N_i=N_n=512$): weights = 9.4 MB $>$ 4.5 MB L2 $\rightarrow$ must go to DRAM.

2.5 K4: regblock (Register Blocking)

Each thread computes `REG_N = 4` consecutive output channels, reusing each input value $4\times$ before discarding:

$$\text{DRAM}_{\text{regblock}} = \frac{\text{input_naive}}{N_{\text{reg}}} + \text{weights} + \text{output}$$

2.6 K5: warp (Warp-Level ni Reduction)

32 threads cooperate on one output element, partitioning the dot product across lanes. Cache-line granularity:

$$\text{input_lines} = B \cdot N_y \cdot N_x \cdot N_n \cdot \left\lceil \frac{K_y K_x N_i}{\text{warp_size}} \right\rceil, \quad \text{DRAM}_{\text{warp}} = 128 \times \text{input_lines} + \text{weights} + \text{output}$$

Part 2b: Occupancy and Coalescing Model for K3 smem_wi

The K3 smem_wi kernel exposes two additional hardware effects that cannot be captured by DRAM traffic alone: **SM occupancy** (limited by shared memory per block) and **cache-line coalescing efficiency** (determined by how many channels are loaded per iteration).

Occupancy Model

Shared memory per block grows linearly with smwi_tile_ni:

$$\text{smem/block} = \text{tni} \times \underbrace{(K_y K_x + (T_y + K_y - 1)(T_x + K_x - 1))}_{=333} \times 4 \text{ bytes} = \text{tni} \times 1332 \text{ bytes}$$

Blocks per SM are limited by whichever resource saturates first:

$$\text{blocks/SM} = \min \left(\left\lfloor \frac{96 \text{ KB}}{\text{smem/block}} \right\rfloor, \frac{2048 \text{ threads/SM}}{256 \text{ threads/block}}, 32 \right)$$

smwi_tile_ni	smem/block	max blocks (smem)	max blocks (threads)	blocks/SM	occupancy
2	2,664 B	36	8	8	1.00
4	5,328 B	18	8	8	1.00
8	10,656 B	9	8	8	1.00
16	21,312 B	4	8	4	0.50
32	42,624 B	2	8	2	0.25

Low occupancy hurts latency hiding: fewer concurrent warps are available to overlap memory stalls with useful computation.

Coalescing Model

A 32-thread warp splits into $\lfloor 32 / \text{tni} \rfloor$ sub-groups, each loading tni consecutive floats from the same (l_y, l_x) position. A 128-byte cache line holds **32 floats**; each sub-group uses $\min(\text{tni}, 32)$ of those 32 slots:

$$\text{coalesce_efficiency} = \frac{\min(\text{tni}, 32)}{32}$$

tni	floats/sub-group	cache-line utilization
2	2	6.25%
4	4	12.5%
8	8	25%
16	16	50%
32	32	100%

Poor coalescing means each cache line fetch contains mostly wasted bytes, reducing effective DRAM bandwidth. L1 caching partially mitigates this (nearby warps share cache lines), captured by a dampened exponent (0.6) rather than the full inverse-efficiency penalty.

Combined Tile-Size Prediction

Starting from the fitted model at ref_tni = 8, corrections are applied multiplicatively to the ideal compute and memory times:

$$t_{\text{compute}}(\text{tni}) = t_{\text{compute,ref}} \times \sqrt{\frac{\text{occ}_{\text{ref}}}{\text{occ}(\text{tni})}} \quad (\text{sqrt: BW-saturated kernels partially mask occupancy loss})$$
$$t_{\text{mem}}(\text{tni}) = t_{\text{mem,ref}} \times \left(\frac{\text{coalesce_eff}_{\text{ref}}}{\text{coalesce_eff}(\text{tni})} \right)^{0.6} \quad (\text{exponent} < 1: \text{L1 re-use dampens raw cache-line waste})$$

$$\text{predicted}(\text{tni}) = \max(s_c \cdot t_{\text{compute}}(\text{tni}), t_{L2}, s_m \cdot t_{\text{mem}}(\text{tni}))$$

where s_c, s_m are the same fitted scale factors as the base model and t_{L2} is the time to serve L2-resident weight reuse (typically small, sub-ms for Conv1).

Validation result (Conv1 224×224 smem_wi sweep):

smwi_tile_ni	Measured (ms)	Predicted (ms)	MAPE	Bottleneck
2	154.9	151.7	2.1%	memory (coalescing)
4	94.1	100.1	6.4%	memory (coalescing)
8	57.1	66.0	15.7%	memory
16	56.8	67.4	18.7%	compute (occ = 0.5)
32	105.6	95.3	9.7%	compute (occ = 0.25)

Average tile-sweep MAPE: **10.5%** vs roofline baseline **95.1%**.

The tni=8/16 over-prediction (~16–18%) occurs because at those tile sizes the kernel sits near the compute/memory crossover, and the occupancy penalty at tni=16 is partially cancelled by improved coalescing — an interaction the simple sqrt/power model does not fully capture.

Part 3: L2 Pressure and Traffic Correction

3.1 L2 Cache Pressure Metric

$$\text{L2_pressure} = \frac{\max(\text{weight_bytes}, \text{input_bytes}, \text{output_bytes})}{\text{L2_size}}$$

TITAN V L2 = 4.5 MB.

Config	Dominant tensor	L2 pressure	Effect
Conv1 naive	input $\approx$ 200 MB	$\approx$ 44	No L2 reuse for input
Conv2 naive	input $\approx$ 8.2 MB	$\approx$ 1.8	Partial L2 reuse
Conv1 smem_wi	input patches $\approx$ 15 GB	$\gg$ 1	Input goes to DRAM; weights stay in L2

3.2 Log-Linear Correction

The analytic DRAM estimate may under- or over-estimate real traffic due to conflict misses, coalescing gaps, and prefetch inefficiency. We fit a **per-kernel log-linear correction**:

$$\text{DRAM_corrected} = \text{DRAM_analytic} \times \exp(a + b \cdot \log(\text{L2_pressure}))$$

Fitting (per kernel, using Conv1 and Conv2 measurements):

1. Compute DRAM_analytic for each sample.
2. Compute $r_i = \text{DRAM_measured}_i / \text{DRAM_analytic}_i$.
3. Fit $\log r = a + b \log(\text{L2_pressure})$ by linear regression.

Correction is clamped to $[0.05, 20]$ to prevent extrapolation blow-up. When the correction factor is < 1 , the analytic model over-estimated traffic and some data was served from L2, not DRAM.

Part 4: Hierarchical Bottleneck Analysis

4.1 Three-Level Ideal Time Decomposition

For each prediction we compute three ideal times:

$$t_{\text{compute}} = \frac{\text{FLOPs}}{\text{peak_FP32}} \times 10^{-9}$$

$$t_{\text{DRAM}} = \frac{\text{DRAM_bytes_corrected}}{\text{DRAM_BW}} \times 10^{-9}$$

$$t_{\text{L2}} = \frac{\text{L2_bytes}}{\text{L2_BW}} \times 10^{-9}$$

L2 bytes are estimated in two cases:

- **Case A** (correction < 1): some analytic traffic was served from L2, not DRAM. `L2_bytes = max(0, DRAM_analytic - DRAM_corrected)`
- **Case B** (weight tensor fits in L2): blocks re-read weights from L2, not DRAM. `L2_bytes = max(0, total_weight_analytic_loads - syn_bytes)`

For Conv1 smem_wi: weights = 144 KB << L2 → L2 serves all inter-block weight reuse ($t_{\text{L2}} \approx 5.3$ ms). For Conv2 smem_wi: weights = 9.4 MB > L2 → L2 bytes ≈ 0.

4.2 Per-Kernel Scaling Factors

Real hardware rarely achieves ideal times. Scale factors s_c (compute) and s_m (memory) absorb stalls, occupancy losses, and other microarchitectural overheads. L2 is assumed to run at rated bandwidth ($s_{\text{L2}} = 1.0$, unscaled) because L2 latency/bandwidth is relatively predictable.

$$\text{predicted_time} = \max(s_c \cdot t_{\text{compute}}, t_{\text{L2}}, s_m \cdot t_{\text{DRAM}})$$

The fitting minimizes MAPE via a grid search over $s_c, s_m \in [0.5, 128]$. Both the fitting loop and the prediction call use the same three-term `max()` so the calibration is consistent.

Bottleneck identification: the model reports which term dominates — "compute", "l2", or "memory" — enabling kernel-specific optimization insight.

Config	Bottleneck (primary model)
Conv1 naive / smem_w / unroll	memory (DRAM)
Conv1 smem_wi	memory (DRAM; weights in L2)
Conv2 smem_wi	compute
Conv2 regblock	compute
Conv2 warp	memory (DRAM)

Part 5: Validation Methodology

5.1 Dataset

We validate on two convolution configurations:

1. **Conv1:** 224×224 input, 64→64 channels, $K_y = K_x = 3$, $B = 16$.
2. **Conv2:** 14×14 input, 512→512 channels, $K_y = K_x = 3$, $B = 16$.

Measured data: kernel runtime (ms) from GPU CUDA events, DRAM bytes (read + write) from Nsight Compute `dram__bytes_read.sum / dram__bytes_write.sum`.

5.2 Two-Path Evaluation

Model	DRAM source	Can generalize?
Shortcut baseline	Measured DRAM bytes	No — needs ncu per config
Primary model	Corrected analytic bytes	Yes — purely analytic

5.3 Error Metrics

$$\text{MAPE} = \frac{|\text{predicted} - \text{measured}|}{\text{measured}} \times 100\%$$

Part 6: Results and Interpretation

6.1 Primary Model Performance

Config	Kernel	Measured (ms)	Predicted (ms)	Time MAPE	Measured GFLOPS	Predicted GFLOPS	GFLOPS MAPE
Conv1 224×224, Ni=Nn=64	K0 naive	686.4	644.4	6.1%	86.2	91.9	6.5%
	K1 smem_w	665.7	625.7	6.0%	88.9	94.6	6.4%
	K2 unroll	1128.6	1058.4	6.2%	52.4	55.9	6.7%
	K3 smem_wi	57.0	66.0	15.8%	1037.6	896.3	13.6%
	K4 regblock	397.2	352.8	11.2%	149.0	167.8	12.6%
	K5 warp	171.0	190.7	11.5%	346.1	310.4	10.3%
Conv2 14×14, Ni=Nn=512	K0 naive	84.2	95.3	13.2%	175.7	155.2	11.6%
	K1 smem_w	99.2	95.3	3.9%	149.2	155.2	4.0%
	K2 unroll	86.9	95.3	9.8%	170.3	155.2	8.9%
	K3 smem_wi	12.5	11.9	4.9%	1181.5	1241.7	5.1%
	K4 regblock	23.4	23.8	1.7%	631.5	620.8	1.7%
	K5 warp	39.3	47.7	21.4%	376.8	310.4	17.6%
Average				9.30%			8.75%

6.2 Comparison with Roofline

Config / Kernel	Roofline MAPE	Primary model MAPE
Conv1 naive	~99.4%	6.1%
Conv2 smem_wi	~97.6%	4.9%
All kernels avg	~97%	9.3%

The model reduces roofline error by ~10× on average by accounting for actual DRAM traffic patterns and L2 cache pressure.

Part 6b: Sweep Validation

Tile Size Sweep (smwi_tile_ni: 2–32)

K3 smem_wi is compiled with different smwi_tile_ni values (conv_tni_* binaries) and measured on Conv1.

smwi_tile_ni	Measured (ms)	Model (ms)	Model MAPE	Roofline (ms)	Roofline MAPE
2	154.9	151.7	2.1%	3.97	97.4%
4	94.1	100.1	6.4%	3.97	95.8%
8	57.1	66.0	15.7%	3.97	93.1%
16	56.8	67.4	18.7%	3.97	93.0%
32	105.6	95.3	9.7%	3.97	96.2%
Average			10.5%		95.1%

The model correctly identifies the U-shaped performance curve: small tni is slow due to coalescing waste (memory bottleneck); large tni is slow due to smem-limited occupancy (compute bottleneck). The sweet spot tni ∈ {8, 16} achieves the lowest runtime.

Problem Size Sweep (Ny = Nx ∈ {16, 32, 56, 112, 224})

Ny = Nx	Measured (ms)	Model (ms)	Model MAPE	Roofline (ms)	Roofline MAPE
16	0.253	0.243	3.9%	0.020	92.0%
32	0.892	0.973	9.1%	0.081	90.9%
56	3.157	2.979	5.6%	0.248	92.1%
112	10.815	11.917	10.2%	0.993	90.8%
224	57.071	66.039	15.7%	3.972	93.0%
Average			8.9%		91.8%

The model scales correctly with Ny² (proportional to input tensor size). The bottleneck transitions from compute (small Ny) to memory (large Ny) as the input patch workload grows — correctly identified by the three-level bottleneck analysis.

Part 7: Generalization to New Kernel Shapes (KY=5, KX=5)

With the analytic DRAM path and L2-pressure correction, the model generalizes to:

- **Different kernel sizes:** KY=5, KX=5 (weight tensor changes size, L2 fit changes).
- **Arbitrary spatial sizes:** any Ny, Nx, Ni, Nn.

Predicted MAPE for KY=KX=5 is expected to be 10–20% (slightly higher due to changed kernel geometry and smem sizes not re-fitted).

Part 8: Architecture Insight — Sensitivity Analysis

A sensitivity analysis sweeps each hardware parameter from 0.25× to 4× of its current value and observes the change in predicted runtime for K3 smem_wi.

Methodology: Hold all other parameters fixed, vary one parameter, recompute:

predicted = max(s_c · t_compute, t_L2, s_m · t_DRAM)

Conv1 (224×224, Ni=Nn=64) — memory-bound

Parameter	Effect of 2× increase	Interpretation
DRAM Bandwidth	~2× speedup	Critical — Conv1 is DRAM-dominated
FP32 TFLOPS	minimal	Not compute-bound even at 4×
L2 Cache Size	moderate below threshold	Weights (144 KB) fit in current L2; only matters if L2 shrinks below that

Conv2 (14×14, Ni=Nn=512) — compute-bound

Parameter	Effect of 2× increase	Interpretation
DRAM Bandwidth	minimal	Already compute-bound
FP32 TFLOPS	~2× speedup	Critical — Conv2 is compute-dominated
L2 Cache Size	minimal	Weights (9.4 MB) exceed L2; larger L2 doesn’t help much

Key architectural recommendation: For large-spatial small-channel layers (Conv1-style), future GPUs should prioritize **DRAM bandwidth** (wider HBM interfaces). For small-spatial large-channel layers (Conv2-style, the dominant ML workload), **tensor core TFLOPS** is the bottleneck — wider/deeper tensor cores help most.

Part 9: Code Organization

9.1 hierarchical1.py (Main Model)

Function	Purpose
conv_flops()	Compute total FLOPs

Function	Purpose
conv_tensor_bytes()	Weight / input / output byte counts
conv_l2_pressure()	L2 pressure metric
_analytic_dram_bytes()	Per-kernel analytic DRAM traffic
_hierarchical_ideal_times_ms()	Three-level ideal times: compute, L2, DRAM
_fit_traffic_correction()	Log-linear DRAM correction per kernel
_fit_kernel_scalers()	Grid-search s_c, s_m using the three-level $\max()$
predict_sample()	Predict time + bottleneck label for a sample
evaluate()	Full pipeline (load, fit, predict)
compute_smwi_occupancy()	SM occupancy vs smwi_tile_ni
smwi_coalesce_efficiency()	Cache-line utilization = $\min(tni, 32)/32$
predict_smwi_tile_sweep()	Tile-size sweep with occupancy + coalescing corrections
predict_smwi_problem_sweep()	Problem-size sweep (analytic model, no ncu needed)

9.2 hw_params.py

Tunable hardware constants for NVIDIA TITAN V (CC 7.0):

Parameter	Value	Source
fp32_tflops	14.9	80 SM $\times$ 64 cores $\times$ 2 $\times$ 1.455 GHz
dram_bw_gb_s	652.8	HBM2, cross-checked with ncu
l2_bw_gb_s	3100	Estimated, tunable
l1_bw_gb_s	12800	Estimated, tunable
l2_bytes	4,718,592	deviceQuery (4.5 MB)
l1_spad_bytes_per_sm	98,304	96 KB unified L1 + shared

9.3 validate.py

Function	Output
plot_model_vs_measured()	Bar chart: measured vs predicted runtime
plot_mape_by_kernel()	MAPE breakdown by kernel
plot_error_vs_tile_size()	Tile sweep: measured / model / roofline + MAPE bars
plot_error_vs_problem_size()	Problem sweep: measured / model / roofline + MAPE bars
plot_sensitivity_analysis()	HW parameter sensitivity for smem_wi

Supports custom result directories via `--results-dir` and `--config`.

9.4 analysis/plot_roofline_predicted.py

Renders a roofline chart using model-predicted AI and GFLOPS (not measured data).

Part 10: Model Characteristics

10.1 What Worked Well

- Three-level hierarchical bottleneck:** including L2 as an active term in $\max(s_c \cdot t_c, t_{L2}, s_m \cdot t_{DRAM})$ makes the model architecturally complete. For Conv1 smem_wi, the L2 term captures weight reuse ($t_{L2} \approx 5.3$ ms), which becomes visible when DRAM bandwidth improves.
- L2 traffic estimation for both correction directions:** the model correctly handles both cases — when the correction factor < 1 (analytic over-estimated, difference served from L2) and when the weight tensor fits in L2 (L2 serves repeated weight accesses across spatial tiles).
- Occupancy + coalescing model for tile sweep:** correctly identifies the U-shaped performance curve: coalescing waste dominates at small tni (memory bottleneck), smem-limited occupancy dominates at large tni (compute bottleneck). Achieves 10.5% average MAPE vs 95.1% for roofline.

- 4. **Correct coalescing formula:** each 128-byte cache line holds 32 floats; coalescing efficiency = $\min(\text{tni}, 32)/32$, not $\min(\text{tni}, 16)/16$. This ensures predictions for tni=32 (100% coalescing) differ correctly from tni=16 (50%).
- 5. **Bottleneck identification:** predict_sample() now returns a "bottleneck" key ("compute", "l2", or "memory"), enabling automatic identification of the limiting resource for any configuration.
- 6. **L2-pressure-aware correction:** log-linear, shape-aware correction effectively bridges the gap between ideal analytic traffic and real hardware behavior.

Appendix A: Hardware Parameters (NVIDIA TITAN V, Compute Capability 7.0)

Parameter	Value	Note
FP32 peak	14.9 TFLOPS	80 SMs × 64 cores × 2 × 1.455 GHz
DRAM bandwidth	652.8 GB/s	HBM2, cross-checked vs ncu at 93.9% utilisation
L2 cache size	4.5 MB	From deviceQuery
L2 bandwidth	3,100 GB/s	Estimated from Volta whitepaper
L1/shared size	96 KB/SM	Unified, configurable
L1 bandwidth	12,800 GB/s	Estimated
Memory clock	850 MHz	From deviceQuery
Warp size	32	
Max threads/SM	2,048	
Max blocks/SM	32	
Num SMs	80	

Appendix B: Roofline Chart Interpretation

- **Left of ridge point** (AI < 22.82 FLOPs/byte): memory-bound → improve bandwidth or reduce traffic.
- **Right of ridge point:** compute-bound → only more FP32 TFLOPS helps.
- **Ridge point:** AI = 14,900 GFLOPS / 652.8 GB/s ≈ 22.82 FLOPs/byte for TITAN V.


Example placements:

- Conv1 naive: AI ≈ 0.14 (far left, memory-bound) → 86 GFLOPS achieved.
- Conv2 smem_wi: AI ≈ 52.9 (right of ridge, compute-bound) → 1181 GFLOPS achieved.

Figures


Model vs Measured Runtime (All Kernels)

Figure: Measured vs Predicted Runtime (log scale) for all 12 kernel/config combinations. Demonstrates the model’s accuracy across every kernel variant for both Conv1 and Conv2.



MAPE by Kernel


Figure: Average MAPE (%) per kernel type, averaged over Conv1 and Conv2. Shows which kernels the model predicts most and least accurately.




Model vs Roofline Comparison

Figure: Roofline vs Hierarchical Model accuracy comparison — top panel shows absolute runtimes (measured / roofline / hierarchical) per kernel; bottom panel shows MAPE bars. The hierarchical model reduces average MAPE from ~97.5% (roofline) to ~9.3%.

Measured vs Roofline vs Hierarchical Model — all kernels




Roofline MAPE vs Hierarchical Model MAPE





Error vs Tile Size

Figure: Error vs Tile Size — Tile sweep for the smem_wi kernel (smwi_tile_ni ∈ {2,4,8,16,32}). This figure illustrates the U-shaped performance curve: poor coalescing at small tile sizes and occupancy limits at large tile sizes; the model captures the sweet spot.




Error vs Problem Size

Figure: Error vs Problem Size — Model vs Measured runtime error across problem sizes ($N_y = N_x \in \{16, 32, 56, 112, 224\}$). The plot shows model error (MAPE) compared to the roofline baseline and highlights that the mechanistic model maintains low error as problem size grows.



Predicted Roofline (Model)


Figure: Roofline predicted by the mechanistic model — AI vs predicted GFLOPS. This roofline uses model-predicted arithmetic intensity and predicted throughput (not measured), showing how kernels move relative to the compute/memory ridge under the model.



Predicted vs Measured DRAM Traffic

Figure: Side-by-side comparison of ncu-measured DRAM bytes vs analytically predicted DRAM bytes (per kernel). MAPE labels on top of each bar pair. Average DRAM MAPE: ~5.4%.

Predicted vs Measured DRAM Traffic — all kernels



Architecture Sensitivity Analysis

Figure (DRAM Bandwidth): Predicted runtime vs DRAM bandwidth multiplier ($0.25\times - 4\times$) for K3 smem_wi on Conv1 and Conv2. Conv1 is strongly bandwidth-sensitive; Conv2 is not.

Sensitivity: DRAM Bandwidth | K3 smem_wi




Figure (FP32 TFLOPS): Predicted runtime vs FP32 peak TFLOPS multiplier. Conv2 is compute-bound and benefits strongly from more TFLOPS; Conv1 is insensitive.

Sensitivity: FP32 Peak TFLOPS | K3 smem_wi





Figure (L2 Cache Size): Predicted runtime vs L2 cache size multiplier. Conv1 with small weights (144 KB) only sees an effect when L2 shrinks below the weight footprint.

Sensitivity: L2 Cache Size | K3 smem_wi



Appendix C: Reproduction Instructions

```
cd ~/CS259_Projects/project2

# 1. Run sweeps to collect measured data
./run_sweep.sh           # generates results/measured/*.csv

# 2. Fit model and print calibrated scale factors
```



```
python3 hierarchical.py
```

```
# 3. Generate all validation plots
```

```
python3 validate.py --config default
```

```
# → results_default/: model_vs_measured_time.png, mape_by_kernel.png,
```

```
# error_vs_tile_size.png, error_vs_problem_size.png,
```

```
# roofline_predicted.png, sensitivity_*.png
```

```
# 4. Validate on different kernel geometry (KY=5, KX=5)
```

```
python3 validate.py --config ky_5_kx_5 \
```

```
--results-dir <path_to_ky5_results> \
```

```
--ky 5 --kx 5
```